SQL OF THE DAY

count people, not rows

07 JUL 2026

StrataScratch · #9942

Nivan RS

THE PROBLEM

Find the Olympics with the most distinct athletes. Return the edition name and the number of distinct athletes.

ties are plausible...

Medium · PostgreSQL

THE SOLUTION

```sql
WITH Olymp AS (
    SELECT games,
           COUNT(DISTINCT id) AS athletes
    FROM olympics_athletes_events
    GROUP BY games
)
SELECT *
FROM Olymp
WHERE athletes = (
    SELECT MAX(athletes) FROM Olymp
);
```

HOW IT WORKS

COUNT DISTINCT

COUNT(DISTINCT id) GROUP BY games counts unique athletes per edition. One person in three events is still one athlete.

FIND THE MAX

The subquery SELECT MAX(athletes) returns the single highest count across all editions.

MATCH THE PEAK

WHERE athletes = (SELECT MAX(...)) keeps every edition at that peak, ties included.

THE LOGIC

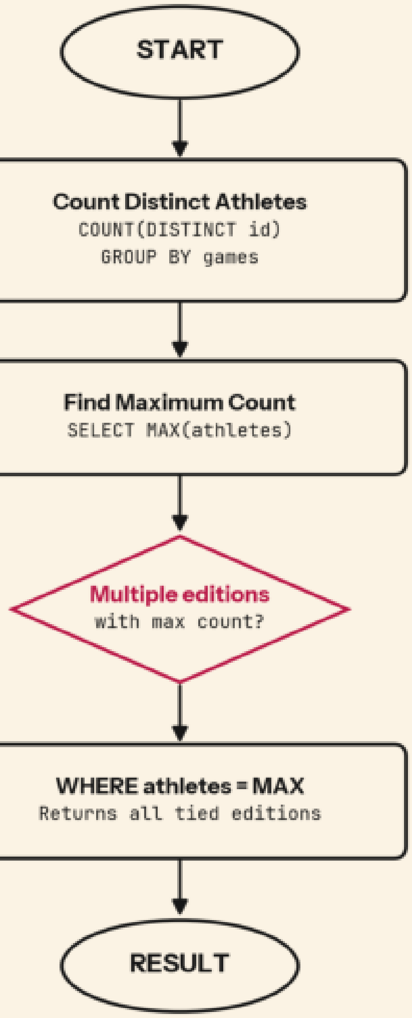

Why not ORDER BY ... LIMIT 1?

LIMIT 1 drops ties with no error.

= MAX() returns every tied edition.

WHY IT MATTERS

01 **Records get quoted**
Most-attended Games lands in media, sponsor decks, and planning models. If two editions tie and you publish one, the record is a coin flip.

02 **Capacity planning**
Host cities size venues and budgets off historical highs, so missing a tied peak understates the real ceiling.

03 **Trend analysis**
Comparing the largest edition across decades only holds up if every tied peak is counted.

KEY TAKEAWAYS

1. Use COUNT(DISTINCT id) when one entity spans many rows. Counting rows overcounts people.
2. Prefer = (SELECT MAX(...)) over ORDER BY + LIMIT 1 when ties can exist.
3. LIMIT 1 is not a bug detector. It returns clean output even while dropping valid data.

follow for more SQL problems

questions? nivanrs@gmail.com

Nivan RS