



SQL OF THE DAY

clean before you split

28 JUL 2026

StrataScratch · #9817

Nivan RS





THE PROBLEM

Find the number of times each word appears across all drafts in `google_file_store`. Output each word and its count, ordered by frequency descending.

punctuation fragments words...

Medium · PostgreSQL





THE SOLUTION

```
WITH raw AS (  
  SELECT LOWER(  
    REGEXP_REPLACE(  
      word, '^[a-zA-Z0-9]', '', 'g'  
    )  
  ) AS word  
FROM google_file_store,  
  UNNEST(  
    STRING_TO_ARRAY(contents, ' ')  
  ) AS word  
)  
SELECT word, COUNT(*) AS occurrences  
FROM raw  
GROUP BY word  
ORDER BY 2 DESC;
```





HOW IT WORKS

SPLIT AND EXPLODE

STRING_TO_ARRAY splits each draft on spaces; UNNEST turns the array into one row per token.

CLEAN EACH TOKEN


REGEXP_REPLACE strips punctuation and LOWER normalizes case, so 'Word.' and 'word' count as one.

COUNT AND SORT

COUNT(*) GROUP BY word aggregates the cleaned tokens, ordered by frequency.



THE LOGIC





WHY IT MATTERS

01

Voice of the corpus

Word frequency surfaces the vocabulary a team actually uses, exposing themes that dominate the drafts rather than the ones leadership assumes are there.

02

Cheap signal before ML

A corpus where 'blocked', 'issue', or 'pending' ranks near the top flags risk and sentiment before anyone builds a model.

03

Document health

Drafts heavy with hedge words like 'maybe' and 'pending' point to unresolved decisions, not finished work.





KEY TAKEAWAYS

- 1 Clean before you split. Punctuation that runs into the next word with no space survives the split, so normalize deliberately.
- 2 LOWER() and REGEXP_REPLACE() together are the minimum normalization for any word-frequency task. Either alone still fragments counts.
- 3 Word frequency is a proxy for attention. What people write most is what they are actually thinking about.

follow for more SQL problems

questions? nivanrs@gmail.com

